



CUDA Performance Considerations (1 of 2)

Patrick Cozzi
University of Pennsylvania
CIS 565 - Spring 2012

Agenda

- Parallel Reduction Revisited
- Warp Partitioning
- Memory Coalescing
- Dynamic Partitioning of SM Resources
- Data Prefetching

Efficient data-parallel algorithms

+

Optimizations based on GPU Architecture

=


Maximum Performance

Parallel Reduction


- Recall *Parallel Reduction* (sum)

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---


Parallel Reduction



Parallel Reduction




Parallel Reduction



Parallel Reduction

- Similar to brackets for a basketball tournament
- $\log(n)$ passes for n elements
- How would you implement this in CUDA?



```

__shared__ float partialSum[];
// ... load into shared memory
unsigned int t = threadIdx.x;
for (unsigned int stride = 1;
     stride < blockDim.x;
     stride *= 2)
{
    __syncthreads();
    if (t % (2 * stride) == 0)
        partialSum[t] +=
            partialSum[t + stride];
}

```

Code from <http://courses.engr.illinois.edu/ece498/al/Syllabus.html>

```

__shared__ float partialSum[];
// ... load into shared memory
unsigned int t = threadIdx.x;
for (unsigned int stride = 1;
     stride < blockDim.x;
     stride *= 2)
{
    __syncthreads();
    if (t % (2 * stride) == 0)
        partialSum[t] +=
            partialSum[t + stride];
}

```

Computing the sum for the
elements in shared memory


Code from <http://courses.engr.illinois.edu/ece498/al/Syllabus.html>

```

__shared__ float partialSum[];
// ... load into shared memory
unsigned int t = threadIdx.x;
for (unsigned int stride = 1;
     stride < blockDim.x;
     stride *= 2)
{
    __syncthreads();
    if (t % (2 * stride) == 0)
        partialSum[t] +=
            partialSum[t + stride];
}

```

Stride:
1, 2, 4, ...



Code from <http://courses.engr.illinois.edu/ece498/al/Syllabus.html>

```

__shared__ float partialSum[];
// ... load into shared memory
unsigned int t = threadIdx.x;
for (unsigned int stride = 1;
     stride < blockDim.x;
     stride *= 2)
{
    __syncthreads();
    if (t % (2 * stride) == 0)
        partialSum[t] +=
            partialSum[t + stride];
}

```

Why?

Code from <http://courses.engr.illinois.edu/ece498/al/Syllabus.html>

```

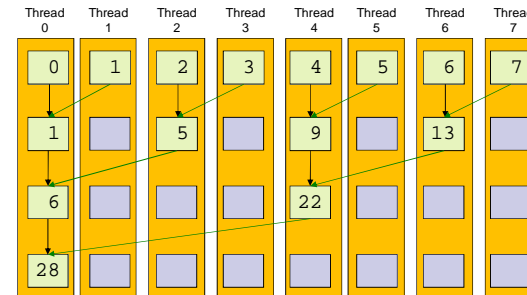
__shared__ float partialSum[];
// ... load into shared memory
unsigned int t = threadIdx.x;
for (unsigned int stride = 1;
     stride < blockDim.x;
     stride *= 2)
{
    __syncthreads();
    if (t % (2 * stride) == 0)
        partialSum[t] +=
            partialSum[t + stride];
}

```

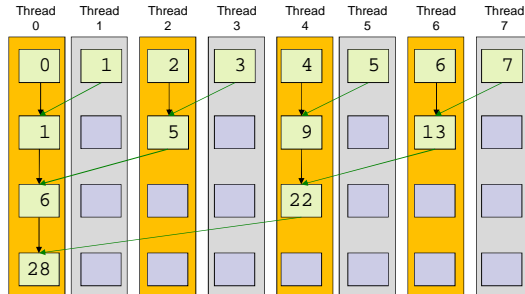
- Compute sum in same shared memory
- As stride increases, what do more threads do?

Code from <http://courses.engr.illinois.edu/ece498/al/Syllabus.html>

Parallel Reduction

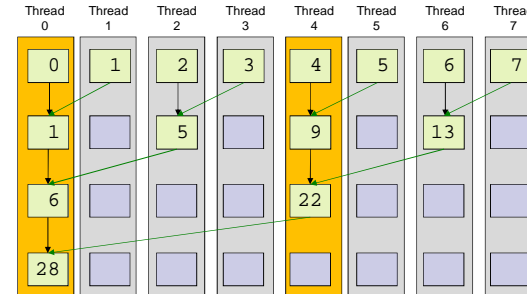


Parallel Reduction



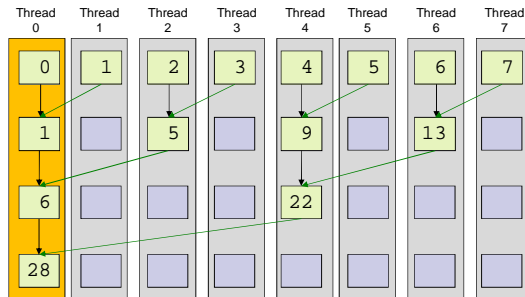
- 1st pass: threads 1, 3, 5, and 7 don't do anything
- Really only need $n/2$ threads for n elements

Parallel Reduction




- 2nd pass: threads 2 and 6 also don't do anything

Parallel Reduction



- 3rd pass: thread 4 also doesn't do anything


Parallel Reduction




- In general, number of required threads cuts in half after each pass

Parallel Reduction


- What if we *tweaked* the implementation?




Parallel Reduction



Parallel Reduction



Parallel Reduction



```
__shared__ float partialSum[]
// ... load into shared memory
unsigned int t = threadIdx.x;
for(unsigned int stride = blockDim.x / 2;
    stride > 0;
    stride /= 2)
{
    __syncthreads();
    if (t < stride)
        partialSum[t] +=
            partialSum[t + stride];
}
```

stride: ..., 4, 2, 1

Code from <http://courses.engr.illinois.edu/ece498/al/Syllabus.html>

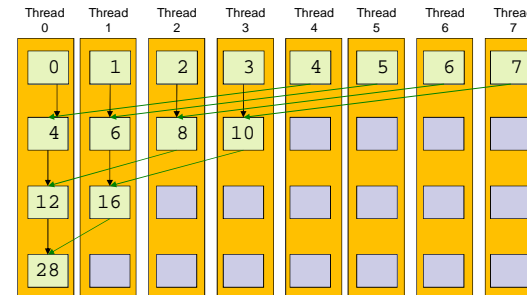
```

__shared__ float partialSum[]
// ... load into shared memory
unsigned int t = threadIdx.x;
for(unsigned int stride = blockDim.x / 2;
    stride > 0;
    stride /= 2)
{
    __syncthreads();
    if (t < stride)
    {
        partialSum[t] +=
            partialSum[t + stride];
    }
}

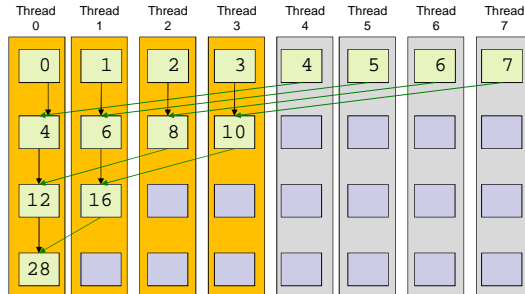
```

Code from <http://courses.engr.illinois.edu/ece498/al/Syllabus.html>

Parallel Reduction

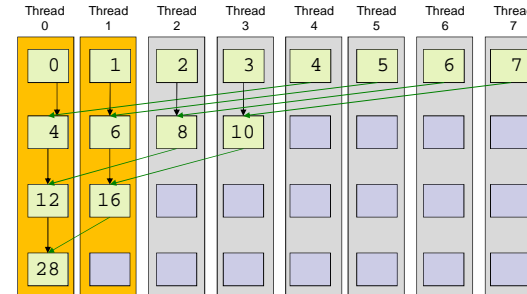


Parallel Reduction




- 1st pass: threads 4, 5, 6, and 7 don't do anything
- Really only need $n/2$ threads for n elements

Parallel Reduction



- 2nd pass: threads 2 and 3 also don't do anything


Parallel Reduction



- 3rd pass: thread 1 also doesn't do anything

Parallel Reduction

- What is the difference?



Parallel Reduction

- What is the difference?

```
if (t % (2 * stride) == 0)
    partialSum[t] +=
        partialSum[t + stride];
```

stride = 1, 2, 4, ...

```
if (t < stride)
    partialSum[t] +=
        partialSum[t + stride];
```

stride = 4, 2, 1, ...

Warp Partitioning

- **Warp Partitioning**: how threads from a block are divided into warps
- Knowledge of warp partitioning can be used to:
 - ☐ Minimize divergent branches
 - ☐ Retire warps early

Understand warp partitioning → make your code run faster

Warp Partitioning

- Partition based on *consecutive increasing* `threadIdx`

Warp Partitioning

■ 1D Block

- `threadIdx.x` between 0 and 512 (G80/GT200)
- Warp n
 - Starts with thread $32n$
 - Ends with thread $32(n + 1) - 1$
- Last warp is padded if block size is not a multiple of 32

Warp 0	Warp 1	Warp 2	Warp 3	...
0...31	32...63	64...95	96...127	...

Warp Partitioning

■ 2D Block

- Increasing `threadIdx` means
 - Increasing `threadIdx.x`
 - Starting with row `threadIdx.y == 0`

Warp Partitioning

■ 2D Block




Image from <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter5-CudaPerformance.pdf>

Warp Partitioning

■ 3D Block

- Start with `threadIdx.z == 0`
- Partition as a 2D block
- Increase `threadIdx.z` and repeat

Warp Partitioning

Divergent branches are within a warp!




Image from: http://bpsi10.idav.ucdavis.edu/talks/03-fatahalian_gpuArchTeraflop_BPS_SIGGRAPH2010.pdf

Warp Partitioning

- For `warpSize == 32`, does any warp have a divergent branch with this code:

```

if (threadIdx.x > 15)
{
    // ...
}
    
```

Warp Partitioning

- For any `warpSize`, does any warp have a divergent branch with this code:

```
if (threadIdx.x > warpSize)
{
    // ...
}
```

Warp Partitioning

- Given knowledge of warp partitioning, which parallel reduction is better?

```
if (t % (2 * stride) == 0)
    partialSum[t] +=
        partialSum[t + stride];
```


stride = 1, 2, 4, ...

```
if (t < stride)
    partialSum[t] +=
        partialSum[t + stride];
```


stride = 4, 2, 1, ...

Warp Partitioning

- Pretend `warpSize == 2`




stride = 1, 2, 4, ...




stride = 4, 2, 1, ...

Warp Partitioning

- 1st Pass




stride = 1, 2, 4, ...



stride = 4, 2, 1, ...


Warp Partitioning

■ 2nd Pass




Warp Partitioning

■ 2nd Pass



Warp Partitioning

■ 2nd Pass



Warp Partitioning

■ Good partitioning also allows warps to be retired early.

- Better hardware utilization

```
if (t % (2 * stride) == 0)
    partialSum[t] +=
    partialSum[t + stride];
```


stride = 1, 2, 4, ...

```
if (t < stride)
    partialSum[t] +=
    partialSum[t + stride];
```

stride = 4, 2, 1, ...


Warp Partitioning

■ Parallel Reduction




Warp Partitioning

■ 1st Pass




Warp Partitioning

■ 1st Pass




Warp Partitioning

■ 2nd Pass



Warp Partitioning

■ 2nd Pass



Memory Coalescing

- Given a matrix stored *row-major* in *global memory*, what is a *thread's* desirable access pattern?




Image from: http://bps10.idav.ucdavis.edu/talks/03-fatahalian_gpuArchTeraflop_BPS_SIGGRAPH2010.pdf

Memory Coalescing

- Given a matrix stored *row-major* in *global memory*, what is a *thread's* desirable access pattern?




Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter5-CudaPerformance.pdf>

Memory Coalescing

- Given a matrix stored *row-major* in *global memory*, what is a *thread's* desirable access pattern?

- ☐ a) column after column
 - *Individual threads* read increasing, consecutive memory address
- ☐ b) row after row
 - *Adjacent threads* read increasing, consecutive memory addresses

Memory Coalescing




Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter5-CudaPerformance.pdf>

Memory Coalescing





Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter5-CudaPerformance.pdf>

Memory Coalescing



Recall warp partitioning; if these threads are in the same warp, global memory addresses are increasing and consecutive across warps.

Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter5-CudaPerformance.pdf>

Memory Coalescing

- Global memory bandwidth (DRAM)
 - G80 – 86.4 GB/s
 - GT200 – 150 GB/s
- Achieve peak bandwidth by requesting large, consecutive locations from DRAM
 - Accessing random location results in much lower bandwidth

Memory Coalescing

- **Memory coalescing** – rearrange access patterns to improve performance
- Useful today but will be less useful with large on-chip caches


Memory Coalescing

- The GPU coalesce consecutive reads in a **half-warp** into a single read
- **Strategy**: read global memory in a coalesce-able fashion into shared memory
 - Then access shared memory randomly at maximum bandwidth
 - Ignoring **bank conflicts** – next lecture

See Appendix G in the NVIDIA CUDA C Programming Guide for coalescing alignment requirements


SM Resource Partitioning

- Recall a SM dynamically partitions resources:




SM Resource Partitioning

- Recall a SM dynamically partitions resources:




SM Resource Partitioning

- We can have
 - 8 blocks of 96 threads
 - 4 blocks of 192 threads
 - But not 8 blocks of 192 threads



SM Resource Partitioning


- We can have (assuming 256 thread blocks)
 - 768 threads (3 blocks) using 10 registers each
 - 512 threads (2 blocks) using 11 registers each



SM Resource Partitioning

- We can have (assuming 256 thread blocks)
 - 768 threads (3 blocks) using 10 registers each
 - 512 threads (2 blocks) using 11 registers each

- More registers decreases thread-level parallelism
 - Can it ever increase performance?



SM Resource Partitioning

- **Performance Cliff:** Increasing resource usage leads to a dramatic reduction in parallelism
 - For example, increasing the number of registers, unless doing so hides latency of global memory access

SM Resource Partitioning

- CUDA Occupancy Calculator

□ http://developer.download.nvidia.com/compute/cuda/CUDA_Occupancy_calculator.xls

Data Prefetching

- Independent instructions between a global memory read and its use can hide memory latency

```
float m = Md[i];  
float f = a * b + c * d;  
float f2 = m * f;
```

Data Prefetching

- Independent instructions between a global memory read and its use can hide memory latency

```
float m = Md[i];  
float f = a * b + c * d;  
float f2 = m * f;
```

Read global memory

Data Prefetching

- Independent instructions between a global memory read and its use can hide memory latency

```
float m = Md[i];  
float f = a * b + c * d;  
float f2 = m * f;
```

Execute instructions
that are not dependent
on memory read

Data Prefetching

- Independent instructions between a global memory read and its use can hide memory latency

```
float m = Md[i];  
float f = a * b + c * d;  
float f2 = m * f;
```

Use global memory after the above line from enough warps hide the memory latency

Data Prefetching

- *Prefetching* data from global memory can effectively increase the number of independent instructions between global memory read and use

Data Prefetching

- Recall tiled matrix multiply:

```
for (/* ... */)   
{  
    // Load current tile into shared memory  
    __syncthreads();  
    // Accumulate dot product  
    __syncthreads();  
}
```

Data Prefetching

- Tiled matrix multiply with prefetch:

```
// Load first tile into registers  
  
for (/* ... */)   
{  
    // Deposit registers into shared memory  
    __syncthreads();  
    // Load next tile into registers  
    // Accumulate dot product  
    __syncthreads();  
}
```

Data Prefetching

■ Tiled matrix multiply with prefetch:

```
// Load first tile into registers

for (/* ... */)
{
    // Deposit registers into shared memory
    __syncthreads();
    // Load next tile into registers
    // Accumulate dot product
    __syncthreads();
}
```

Data Prefetching

■ Tiled matrix multiply with prefetch:

```
// Load first tile into registers

for (/* ... */)
{
    // Deposit registers into shared memory
    __syncthreads();
    // Load next tile into registers
    // Accumulate dot product
    __syncthreads();
}
```

Prefetch for next iteration of the loop

Data Prefetching

■ Tiled matrix multiply with prefetch:

```
// Load first tile into registers

for (/* ... */)
{
    // Deposit registers into shared memory
    __syncthreads();
    // Load next tile into registers
    // Accumulate dot product
    __syncthreads();
}
```

These instructions executed by enough threads will hide the memory latency of the prefetch